

Analysis of Yelp Dataset

I will analyze Yelp dataset. This dataset is downloaded from [Kaggle](#) and uploaded to my S3 bucket.

Part I: Installation and Initial Setup

In [1]: `%%info`

```
Current session configs: {'conf': {'spark.pyspark.python': 'python3', 'spark.pyspark.virtualenv.enabled':  
'true', 'spark.pyspark.virtualenv.type': 'native', 'spark.pyspark.virtualenv.bin.path':  
'/usr/bin/virtualenv'}, 'kind': 'pyspark'}
```

No active sessions.

In [2]: `# check currently installed packages
sc.list_packages()`

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
11	application_1606143194499_0012	pyspark	idle	Link	Link	✓

SparkSession available as 'spark'.

Package	Version
beautifulsoup4	4.9.1
boto	2.49.0
click	7.1.2
jmespath	0.10.0
joblib	0.16.0
lxml	4.5.2
mysqlclient	1.4.2
nltk	3.5
nose	1.3.4
numpy	1.16.5
pip	9.0.1
py-dateutil	2.2
python37-sagemaker-pyspark	1.4.0
pytz	2020.1
PyYAML	5.3.1
regex	2020.7.14

```
setuptools          28.8.0
six                  1.13.0
soupsieve            1.9.5
tqdm                 4.48.2
wheel                0.29.0
windmill             1.6
```

```
In [3]: # install and update necessary dependencies (pandas, matplotlib, seaborn, and numpy)
sc.install_pypi_package("pandas==1.0.3")
sc.install_pypi_package("matplotlib==3.2.1")
sc.install_pypi_package("seaborn==0.10.0")
```

Collecting pandas==1.0.3

Using cached https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1edc0163743f55aaeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)

Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from pandas==1.0.3)

Collecting python-dateutil>=2.6.1 (from pandas==1.0.3)

Using cached https://files.pythonhosted.org/packages/d4/70/d60450c3dd48ef87586924207ae8907090de0b306af2bce5d134d78615cb/python_dateutil-2.8.1-py2.py3-none-any.whl

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3)

Installing collected packages: python-dateutil, pandas

Successfully installed pandas-1.0.3 python-dateutil-2.8.1

Collecting matplotlib==3.2.1

Using cached https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3balf09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1606254383641-0/lib/python3.7/site-packages (from matplotlib==3.2.1)

Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/8a/bb/488841f56197b13700afd5658fc279a2025a39e22449b7cf29864669b15d/pyparsing-2.4.7-py2.py3-none-any.whl>

Collecting cyclor>=0.10 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/f7/d2/e07d3ebb2bd7af696440ce7e754c59dd546ffe1bbe732c8ab68b9c834e61/cyclor-0.10.0-py2.py3-none-any.whl>

Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)

Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)

Using cached https://files.pythonhosted.org/packages/d2/46/231de802ade4225b76b96cffe419cf3ce52bbe92e3b092cf12db7d11c207/kiwisolver-1.3.1-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)

Installing collected packages: pyparsing, cyclor, kiwisolver, matplotlib

Successfully installed cyclor-0.10.0 kiwisolver-1.3.1 matplotlib-3.2.1 pyparsing-2.4.7

Collecting seaborn==0.10.0

Downloading <https://files.pythonhosted.org/packages/70/bd/5e6bf595fe6ee0f257ae49336dd180768c1ed3d7c7155b2fdf8>

```
94c1c808a/seaborn-0.10.0-py3-none-any.whl (215kB)
Requirement already satisfied: pandas>=0.22.0 in /mnt/tmp/1606254383641-0/lib/python3.7/site-packages (from sea
born==0.10.0)
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from seaborn==0.10.0)
Collecting scipy>=1.0.1 (from seaborn==0.10.0)
  Using cached https://files.pythonhosted.org/packages/dc/7e/8f6a79b102calea928bae8998b05bf5dc24a90571db13cd119
f275ba6252/scipy-1.5.4-cp37-cp37m-manylinux1_x86_64.whl
Requirement already satisfied: matplotlib>=2.1.2 in /mnt/tmp/1606254383641-0/lib/python3.7/site-packages (from
seaborn==0.10.0)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas>=0.22.0->sea
born==0.10.0)
Requirement already satisfied: python-dateutil>=2.6.1 in /mnt/tmp/1606254383641-0/lib/python3.7/site-packages
(from pandas>=0.22.0->seaborn==0.10.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /mnt/tmp/1606254383641-0/lib/python
3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
Requirement already satisfied: cycler>=0.10 in /mnt/tmp/1606254383641-0/lib/python3.7/site-packages (from matpl
otlib>=2.1.2->seaborn==0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /mnt/tmp/1606254383641-0/lib/python3.7/site-packages (from
matplotlib>=2.1.2->seaborn==0.10.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1-
>pandas>=0.22.0->seaborn==0.10.0)
Installing collected packages: scipy, seaborn
Successfully installed scipy-1.5.4 seaborn-0.10.0
```

Importing Packages

```
In [4]: import pandas as pd
        from pandas import DataFrame
        import numpy as np
        import matplotlib.pyplot as plt
        import matplotlib
```

Loading Business Data

```
In [5]: business = spark.read.json('s3://sta9760f2020spark-taejun/yelp_academic_dataset_business.json')
```

```
In [6]: print('Data frame type: ' + str(type(business)))
```

```
Data frame type: <class 'pyspark.sql.dataframe.DataFrame'>
```

Overview of Data

```
In [7]: print(f'Total Columns: {len(business.dtypes)}')
        print(f'Total Rows: {business.count():,}')
```

```
Total Columns: 14
Total Rows: 209,393
```

```
In [8]: business.printSchema()
```

```
root
|-- address: string (nullable = true)
|-- attributes: struct (nullable = true)
|   |-- AcceptsInsurance: string (nullable = true)
|   |-- AgesAllowed: string (nullable = true)
|   |-- Alcohol: string (nullable = true)
|   |-- Ambience: string (nullable = true)
|   |-- BYOB: string (nullable = true)
|   |-- BYOBCorkage: string (nullable = true)
|   |-- BestNights: string (nullable = true)
|   |-- BikeParking: string (nullable = true)
|   |-- BusinessAcceptsBitcoin: string (nullable = true)
|   |-- BusinessAcceptsCreditCards: string (nullable = true)
|   |-- BusinessParking: string (nullable = true)
|   |-- ByAppointmentOnly: string (nullable = true)
|   |-- Caters: string (nullable = true)
|   |-- CoatCheck: string (nullable = true)
|   |-- Corkage: string (nullable = true)
|   |-- DietaryRestrictions: string (nullable = true)
|   |-- DogsAllowed: string (nullable = true)
|   |-- DriveThru: string (nullable = true)
|   |-- GoodForDancing: string (nullable = true)
|   |-- GoodForKids: string (nullable = true)
|   |-- GoodForMeal: string (nullable = true)
|   |-- HairSpecializesIn: string (nullable = true)
|   |-- HappyHour: string (nullable = true)
|   |-- HasTV: string (nullable = true)
|   |-- Music: string (nullable = true)
|   |-- NoiseLevel: string (nullable = true)
|   |-- Open24Hours: string (nullable = true)
|   |-- OutdoorSeating: string (nullable = true)
|   |-- RestaurantsAttire: string (nullable = true)
|   |-- RestaurantsCounterService: string (nullable = true)
|   |-- RestaurantsDelivery: string (nullable = true)
|   |-- RestaurantsGoodForGroups: string (nullable = true)
```

```

|         -- RestaurantsPriceRange2: string (nullable = true)
|         -- RestaurantsReservations: string (nullable = true)
|         -- RestaurantsTableService: string (nullable = true)
|         -- RestaurantsTakeOut: string (nullable = true)
|         -- Smoking: string (nullable = true)
|         -- WheelchairAccessible: string (nullable = true)
|         -- WiFi: string (nullable = true)
| -- business_id: string (nullable = true)
| -- categories: string (nullable = true)
| -- city: string (nullable = true)
| -- hours: struct (nullable = true)
|     | -- Friday: string (nullable = true)
|     | -- Monday: string (nullable = true)
|     | -- Saturday: string (nullable = true)
|     | -- Sunday: string (nullable = true)
|     | -- Thursday: string (nullable = true)
|     | -- Tuesday: string (nullable = true)
|     | -- Wednesday: string (nullable = true)
| -- is_open: long (nullable = true)
| -- latitude: double (nullable = true)
| -- longitude: double (nullable = true)
| -- name: string (nullable = true)
| -- postal_code: string (nullable = true)
| -- review_count: long (nullable = true)
| -- stars: double (nullable = true)
| -- state: string (nullable = true)

```

```

In [9]: print('Columns overview')
pd.DataFrame(business.dtypes, columns = ['Column Name', 'Data type'])

```

```

Columns overview
Column Name      Data type
0      address      string
1      attributes  struct<AcceptsInsurance:string,AgesAllowed:str...
2      business_id      string
3      categories      string
4      city      string
5      hours  struct<Friday:string,Monday:string,Saturday:st...
6      is_open      bigint
7      latitude      double
8      longitude      double
9      name      string
10     postal_code      string
11     review_count      bigint
12      stars      double
13      state      string

```

```
In [10]: #Display the first 5 rows with the following columns: business_id, name, city, state, categories
business.select('business_id','name','city','state','categories').show(5)
```

business_id	name	city	state	categories
f9NumwFMBDn751xgF...	The Range At Lake...	Cornelius	NC	Active Life, Gun/...
YzvJg0SayhoZgCljU...	Carlos Santo, NMD	Scottsdale	AZ	Health & Medical,...
XNoUzKckATkOD1hP6...	Felinus	Montreal	QC	Pets, Pet Service...
6OAZjbxqM5ol29BuH...	Nevada House of Hose	North Las Vegas	NV	Hardware Stores, ...
51M2Kk903DFYI6gnB...	USE MY GUY SERVIC...	Mesa	AZ	Home Services, Pl...

only showing top 5 rows

Part II: Analyzing Categories

```
In [11]: # The table described from my original yelp dataframe
business.select('business_id','categories').show(5)
```

business_id	categories
f9NumwFMBDn751xgF...	Active Life, Gun/...
YzvJg0SayhoZgCljU...	Health & Medical,...
XNoUzKckATkOD1hP6...	Pets, Pet Service...
6OAZjbxqM5ol29BuH...	Hardware Stores, ...
51M2Kk903DFYI6gnB...	Home Services, Pl...

only showing top 5 rows

```
In [12]: # Split categories to each distinct category
from pyspark.sql.functions import explode, split

business_id_categories = business.withColumn("categories", explode(split('categories', ", ")))
```

```
In [13]: # Display the first 5 rows of the association table
business_id_categories.select('business_id', "categories").show(5)
```

business_id	categories
-------------	------------

```
+-----+
|f9NumwFMBDn751xgF...|Active Life|
|f9NumwFMBDn751xgF...|Gun/Rifle Ranges|
|f9NumwFMBDn751xgF...|Guns & Ammo|
|f9NumwFMBDn751xgF...|Shopping|
|Yzvjq0SayhoZgCljU...|Health & Medical|
+-----+
only showing top 5 rows
```

Total Unique Categories

```
In [14]: unique_categories = business_id_categories.select('business_id', "categories")
unique_categories.select("categories").distinct().count()
```

1336

Top Categories by Business

```
In [15]: business_id_categories.groupby("categories").count().show(20)
```

```
+-----+
|categories|count|
+-----+
|Dermatologists|341|
|Paddleboarding|36|
|Aerial Tours|28|
|Hobby Shops|828|
|Bubble Tea|720|
|Embassy|13|
|Handyman|682|
|Tanning|938|
|Aerial Fitness|29|
|Tempura|1|
|Falafel|159|
|Outlet Stores|399|
|Summer Camps|318|
|Clothing Rental|55|
|Sporting Goods|2311|
|Cooking Schools|118|
|College Counseling|15|
|Lactation Services|50|
|Ski & Snowboard S...|50|
|Museums|359|
+-----+
only showing top 20 rows
```


Bar Chart of Top Categories

```
In [16]: top_categories = business_id_categories.groupby("categories").count().orderBy('count', ascending=False).limit(20)
top_categories = top_categories.set_index('categories', 'count')
top_categories = top_categories.sort_values(by='count', ascending=True)
```

```
In [17]: top_categories.plot.barh(figsize=(15,10), rot=45)
plt.xlabel('Count')
plt.ylabel('Categories')
plt.title('Top Categories by Business', fontsize=30, fontweight='bold')
```

```
Text(0.5, 1.0, 'Top Categories by Business')
```

```
In [18]: %matplotlib plt
```

Part III: Do Yelp Reviews Skew Negative?

Loading Review Data

```
In [19]: review= spark.read.json('s3://sta9760f2020spark-taejun/yelp_academic_dataset_review.json')
```

```
In [20]: print('Data frame type: ' + str(type(review)))
```

```
Data frame type: <class 'pyspark.sql.dataframe.DataFrame'>
```

```
In [21]: print(f'Total Columns: {len(review.dtypes)}')
print(f'Total Rows: {review.count():,}')

```

```
Total Columns: 9
Total Rows: 8,021,122
```

```
In [22]: review.printSchema()
```

```
root
|-- business_id: string (nullable = true)
|-- cool: long (nullable = true)
|-- date: string (nullable = true)
|-- funny: long (nullable = true)
|-- review_id: string (nullable = true)
|-- stars: double (nullable = true)
|-- text: string (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
```

```
In [23]: # business and stars data
business_id_stars = review.select('business_id', "stars")
business_id_stars.show(5)
```

```
+-----+-----+
|      business_id|stars|
+-----+-----+
|-MhfebM0QIsKt87iD...|  2.0|
|lbrU8StCq3yDfr-QM...|  1.0|
|Hql28KMwrEKHqhFrr...|  5.0|
|5JxlZaqCnk1MnbgRi...|  1.0|
|IS4cv902ykd8wj1TR...|  4.0|
+-----+-----+
only showing top 5 rows
```

```
In [24]: from pyspark.sql.functions import avg
business_id_avg_stars = review.groupby("business_id").avg("stars")
```

```
In [25]: business_id_avg_stars.show(5)
```

```
+-----+-----+
|      business_id|      avg(stars)|
+-----+-----+
|VHsNB3pdGVcRgs6C3...| 3.411764705882353|
|RMjCnixEY5i12Ciqn...| 3.5316455696202533|
|ipFreSFhjClfNETuM...|                2.6|
|dLDMU8bOLnkDTmPur...| 4.942857142857143|
|Qm2datcYBPXrPATVG...| 4.352941176470588|
+-----+-----+
only showing top 5 rows
```

```
In [26]: # Join two dataframes by busienss_id
df1 = business_id_avg_stars.select('business_id', 'avg(stars)')
df2 = business.select('business_id', 'stars', 'name', 'city', 'state')
joined_df = df1.join(df2, df1.business_id == df2.business_id)
```

```
In [27]: # Display 5 first rows
joined_df = joined_df.select('avg(stars)', 'stars', 'name', 'city', 'state')
joined_df.show(5)
```

```
+-----+-----+-----+-----+-----+
|      avg(stars)|stars|      name|      city|state|
+-----+-----+-----+-----+-----+
|4.11784140969163| 4.0|Delmonico Steakhouse|Las Vegas|NV|
|                4.5| 4.5|Mr. Pancho Mexica...|Mesa|AZ|
|                3.75| 4.0|Maricopa County D...|Phoenix|AZ|
|                4.0| 4.0|Double Play Sport...|Las Vegas|NV|
|                2.6875| 2.5|Impressions Dental|Chandler|AZ|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [28]: # Dataframe to Pandas dataframe
skew_df = joined_df.select('avg(stars)', 'stars').toPandas()
```

```
In [29]: skew_df["skew"] = (skew_df['avg(stars)'] - skew_df['stars']) / skew_df['stars']
skew_df
```

	avg(stars)	stars	skew
0	3.411765	3.5	-0.025210
1	3.531646	3.5	0.009042
2	2.600000	2.5	0.040000
3	4.942857	5.0	-0.011429
4	4.352941	4.5	-0.032680
...
209388	4.000000	4.0	0.000000
209389	3.000000	3.0	0.000000
209390	5.000000	5.0	0.000000
209391	1.666667	1.5	0.111111
209392	3.250000	3.0	0.083333


[209393 rows x 3 columns]

Skewed Distribution of Reviews

```
In [30]: import seaborn as sns

plt.figure(figsize=(10,5))
sns.distplot(skew_df["skew"], kde=True)
plt.title('Skewness of Reviews', size = 15)
plt.axis((-0.6, 1.30, 0, 11))

%matplotlib plt
```



Calculating Skewness

```
In [31]: # Calculate Pearson's coefficient of skewness with mean, median, and standard deviation  
# round up values to calaute values easily  
  
mean = skew_df['skew'].mean().round(5)  
median = skew_df['skew'].median().round(5)  
std_dev = skew_df['skew'].std().round(5)  
skewness = (3 * (mean-median) / std_dev).round(5)  
  
print('Mean: ', mean)  
print('Median: ', median)  
print('Standard Deviation: ', std_dev)  
print('Skewness: ', skewness)
```

Mean: 0.00124

```
Median: 0.0
Standard Deviation: 0.05524
Skewness: 0.06734
```

IMPLICATIONS

Once I looked at the above graph, I could not decide whether skewness of Yelp reviews is positive or negative. So, I calculated Pearson's coefficient of skewness with mean, median, and standard deviation.

- The formula to calculate skewness is $3 * (\text{Mean} - \text{Median}) / \text{Standard Deviation}$
- Mean is 0.00124
- Median is 0.0
- Standard Deviation is 0.05524

By the formula and values, I could get skewness of 0.06734. It is greater than 0 so I can say skewness is positive. It shows that users who wrote reviews were more satisfied than normal. However, it has not a strong positive value and is close to 0, so I would like to say that it is slightly more satisfied than normal.

Part IV: Should the Elite be Trusted?

Loading user data

```
In [32]: user= spark.read.json('s3://sta9760f2020spark-taejun/yelp_academic_dataset_user.json')
```

```
In [33]: print('Data frame type: ' + str(type(user)))
```

```
Data frame type: <class 'pyspark.sql.dataframe.DataFrame'>
```

```
In [34]: print(f'Total Columns: {len(user.dtypes)}')
print(f'Total Rows: {user.count():,}')
```

```
Total Columns: 22
Total Rows: 1,968,703
```

```
In [35]: user.printSchema()
```

```

root
|-- average_stars: double (nullable = true)
|-- compliment_cool: long (nullable = true)
|-- compliment_cute: long (nullable = true)
|-- compliment_funny: long (nullable = true)
|-- compliment_hot: long (nullable = true)
|-- compliment_list: long (nullable = true)
|-- compliment_more: long (nullable = true)
|-- compliment_note: long (nullable = true)
|-- compliment_photos: long (nullable = true)
|-- compliment_plain: long (nullable = true)
|-- compliment_profile: long (nullable = true)
|-- compliment_writer: long (nullable = true)
|-- cool: long (nullable = true)
|-- elite: string (nullable = true)
|-- fans: long (nullable = true)
|-- friends: string (nullable = true)
|-- funny: long (nullable = true)
|-- name: string (nullable = true)
|-- review_count: long (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
|-- yelping_since: string (nullable = true)

```

```

In [36]: # See necessary columns
user.select('user_id', 'elite', 'average_stars').show(10)

```

user_id	elite	average_stars
ntlvfPzc8eglgvk92...		3.57
FOBRPlBHa3WPHFB5q...	2008,2009,2010,20...	3.84
zZUnPeh2hEp0WydbA...	2010	3.44
QaELAmRcDc5TfJEyl...	2009	3.08
xvu8G900tezTzbbfq...	2009,2010,2011,20...	4.37
z5_82komKV3mI4ASG...	2007	2.88
ttumcu6hWshk_EJWV...		4.0
f4_MRNHvN-yRn7EA8...	2011,2012,2013,20...	3.63
UYACF30806j2mfbB5...		3.75
QG13XBbgHWydzThRB...	2008,2009	4.1

only showing top 10 rows

```

In [37]: # Select only elite users by filter empty value out
elite = user.filter(user['elite'] != '').select('user_id', 'elite', 'average_stars')

# Select columns from Review dataset

```

```

review_info = review.select('business_id', 'stars', 'user_id')

# Join two dataframes to get the elite's reviews
elite_review = elite.join(review_info, elite.user_id == review_info.user_id).drop(review['user_id'])
elite_review.show(5)

```

user_id	elite	average_stars	business_id	stars
1Dul59QEe-Q-7OQHT...	2017	3.75	-8F04F54iDT6VgWPC...	4.0
3pMczoCBOSKBcqMhV...	2017	4.24	p200k46G_A000nCWL...	5.0
jO44Apni7iJZVVK4H...	2010,2011,2012,2013	3.75	jyFoxS8MofdpkAAK6...	1.0
RO78oDy7vbEcOJU8a...	2017,2018	3.98	ewty6EB70nwPJsUkA...	4.0
TFxeEvpjMNQ3AWL49...	2017,2018	3.89	0M3KCmdY-_xlIu5vE...	5.0

only showing top 5 rows

```

In [38]: # Get the average starts per each business from Part III
business_id_avg_stars.show(5)

```

business_id	avg(stars)
VHsNB3pdGVcRgs6C3...	3.411764705882353
RMjCnixEY5i12Ciqn...	3.5316455696202533
ipFreSFhjClfNETuM...	2.6
dLDMU8bOLnkDTmPur...	4.942857142857143
Qm2datcYBPXrPATVG...	4.352941176470588

only showing top 5 rows

```

In [39]: # Join Elite review and Business average stars dataframes
elite_review_skew = elite_review.join(business_id_avg_stars, business_id_avg_stars.business_id == elite_review.
elite_review_skew.show(5)

```

user_id	elite	average_stars	stars	business_id	avg(stars)
PMz5Pp3g19MjVmUdT...	2017,2018	3.79	5.0	--9e1ONYQuAa-CB_R...	4.11784140969163
y22jM83qCgsNGryv9...	2012,2013,2014,20...	3.67	5.0	--9e1ONYQuAa-CB_R...	4.11784140969163
jOLdczkaKqfaFb_a...	2017,2018	4.66	5.0	--9e1ONYQuAa-CB_R...	4.11784140969163
le3Xn9Mheb1hNxWur...	2010,2011,2012,20...	3.86	5.0	--9e1ONYQuAa-CB_R...	4.11784140969163
ERb2hZxyCN9Wz999z...	2017,2018	4.24	4.0	--9e1ONYQuAa-CB_R...	4.11784140969163

+-----+-----+-----+-----+-----+
only showing top 5 rows


```
In [40]: # Calcualte skew value
         elite_review_skew = elite_review_skew.toPandas()
         elite_review_skew['skew'] = (elite_review_skew['stars'] - elite_review_skew['avg(stars)']) / elite_review_skew[
         elite_review_skew
```

		user_id	...	skew
0	PMz5Pp3g19MjVmUdT6JTew	...	0.214228	
1	y22jM83qCgsNGryv9328gQ	...	0.214228	
2	jOLdczkaKqfaFb_a_0A6CA	...	0.214228	
3	le3Xn9Mheb1hNxWurF7RSw	...	0.214228	
4	ERb2hZxyCN9Wz999z4a7Cg	...	-0.028617	
...	
1756322	EtxsD-Jbyxh7TsaWHtGCew	...	0.250000	
1756323	_BQMCYEIfsn_Ldb4W3TjUg	...	-0.642857	
1756324	MMf0LhEk5tGa1LvN7zcDnA	...	0.071429	
1756325	r_B6CZ-gZNKeW1bOzSAXPw	...	-0.785714	
1756326	pUjccQauq08ICVR480S-Lg	...	-0.700000	

[1756327 rows x 7 columns]

Skewed Distribution of Elite Reviews

```
In [41]: plt.figure(figsize=(10,5))
         sns.distplot(elite_review_skew['skew'], kde = True)
         plt.title('Skewness of Elite Reviews', size = 20)
         plt.axis((-1.5, 2, 0, 2))
         %matplotlib plt
```



Calculating Skewness

```
In [42]: # Calculate Pearson's coefficient of skewness with mean, median, and standard deviation
# round up values to calaute values easily

mean = elite_review_skew['skew'].mean().round(5)
median = elite_review_skew['skew'].median().round(5)
std_dev = elite_review_skew['skew'].std().round(5)
skewness = (3 * (mean-median) / std_dev).round(5)

print('Mean: ', mean)
print('Median: ', median)
print('Standard Deviation: ', std_dev)
print('Skewness: ', skewness)
```

Mean: 0.04378

Median: 0.05655
Standard Deviation: 0.30707
Skewness: -0.12476

IMPLICATIONS

A "Skewness of Elite Reviews" graph did not clearly show if skewness is positive or negative. I calculated Pearson's coefficient of skewness with mean, median, and standard deviation.

- The formula to calculate skewness is $3 * (\text{Mean} - \text{Median}) / \text{Standard Deviation}$
- mean is 0.04378
- median is 0.05655
- standard deviation is 0.30707

As a result, I could get skewness of -0.12476. It is less than 0, so I can say skewness is negative. It shows that Elite users were less satisfied than normal business reviews. They might tend to write critic reviews due to the title of the Elite. It would result in leaving more negative reviews than normal business reviews. Skewness is not a strong negative value and is close to 0. It looks like Elite users do not intentionally leave negative reviews.

Elite users would be trusted. However, we don't need to strongly trust them or consider only their reviews rather than normal users who left reviews.